

Case Study: Player Contact Triage Automation

Role: Automation & AI Specialist **Time budget:** maximum 3 hours of hands-on work (honesty based, we will ask how you spent it) **Follow-up:** 45 minute walkthrough call where you present and we work through changes together

Scenario

Northline Interactive is an online casino operator running brands in several regulated and licensed markets: Ontario (Canada), Malta (MGA), Ireland, India and New Zealand. The player support team receives around 900 inbound messages a day across email and live chat. Today a human agent reads every message, decides what it is about, sets a priority and either answers it or routes it to a specialist team.

Support leadership wants a triage layer in front of the agents: something that reads each incoming message, classifies it, sets priority, and decides whether it can be answered automatically, needs a human agent, or must be escalated to a specialist team. You have been asked to build a working prototype and recommend how it should be rolled out.

Your task

You are given `dataset_player_messages.csv` with 40 real-looking inbound messages. Build a working prototype that processes all 40 and outputs, per message:

1. **Category**, one of: Payments & Withdrawals / KYC & Verification / Bonuses & Promotions / Technical & Game Issues / Responsible Gambling / Fraud & Compliance / Complaints & Regulatory / General.
2. **Priority**: low / medium / high / critical.
3. **Route**: auto-respond / human agent / specialist escalation (name the team).
4. Anything else you believe the triage layer should surface. This is deliberately open. Part of the exercise is what you notice.

Use whatever stack you want: n8n, Make, Zapier, Power Automate, UiPath, Python, an LLM API, a spreadsheet with formulas, or any combination. We care about your judgment, not your tool loyalty. The prototype must actually run; screenshots of an imagined design are not enough.

Constraints (treat these as real)

- Northline operates in **regulated markets**. Responsible gambling contacts carry legal obligations and hard deadlines in some jurisdictions. Getting one wrong is a licence issue, not a support issue.

- **Player data is sensitive.** Assume payment card data and identity documents must not be sent to any external service that is not an approved processor. If you use a hosted LLM API, you must deal with this somehow and explain how.
- Regulators can audit. Anything the automation decides needs to be **explainable and traceable** after the fact.
- Assume a modest budget. If your design calls per-message on a paid LLM API, know roughly what 900 messages a day costs.

Deliverables

1. **The prototype** (repo, export file, or shared workspace) plus the **processed output for all 40 messages** (CSV or JSON).
2. **A decision log, maximum 2 pages:** what you chose to automate fully, what you kept human and why, what failure modes worry you, what you would need before putting this in production, and a rough cost estimate at 900 messages/day.
3. **Your AI transcripts.** We expect and encourage you to use AI assistants (Claude, ChatGPT, Copilot, whatever you like) to build this. Export or screenshot the full conversations and include them. There is no penalty for heavy AI use. We are hiring for exactly this skill and we want to see how you drive these tools.

The walkthrough call

You will demo the prototype live, and we will ask you to change it in front of us based on new requirements we introduce during the call. Make sure you can open and edit whatever you built.

What we are evaluating

- Whether your triage decisions are safe and sensible for a regulated gambling business.
- How you split work between automation, AI and humans.
- How well you use AI tools as a working partner.
- Whether you can defend and adapt your design under new constraints.

There is no single correct architecture. A simple design with sound judgment beats an impressive design with unsafe defaults.